

Untitled

June 18, 2022

```
[1]: ### IMPORT: -----
import scipy.stats as stats
import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
%matplotlib inline
import seaborn as sns
import warnings
warnings.filterwarnings('ignore') # To suppress warnings
# set the background for the graphs
from scipy.stats import skew
plt.style.use('ggplot')
from sklearn.model_selection import train_test_split # Sklearn package's
    ↳ randomized data splitting function
from sklearn.preprocessing import StandardScaler
from sklearn.preprocessing import OneHotEncoder
import math
from sklearn.linear_model import LinearRegression
from sklearn.metrics import mean_absolute_error, mean_squared_error, r2_score
pd.set_option('display.float_format', lambda x: '%.3f' % x)
pd.set_option('display.max_rows', 300)
pd.set_option('display.max_colwidth', 400)
pd.set_option('display.float_format', lambda x: '%.5f' % x) # To suppress
    ↳ numerical display in scientific notations
import statsmodels.api as sm
print("Load Libraries- Done")
```

Load Libraries- Done

```
[2]: # importing dataset
data=pd.read_csv('CAR DETAILS FROM CAR DEKHO.csv')
```

```
[3]: data.head(5)
```

```
[3]:
```

	name	year	selling_price	km_driven	fuel	\
0	Maruti 800 AC	2007	60000	70000	Petrol	
1	Maruti Wagon R LXI Minor	2007	135000	50000	Petrol	

2	Hyundai Verna 1.6 SX	2012	600000	100000	Diesel
3	Datsun RediGO T Option	2017	250000	46000	Petrol
4	Honda Amaze VX i-DTEC	2014	450000	141000	Diesel

	seller_type	transmission	owner
0	Individual	Manual	First Owner
1	Individual	Manual	First Owner
2	Individual	Manual	First Owner
3	Individual	Manual	First Owner
4	Individual	Manual	Second Owner

1 assumptions

- normality: our data is looking to be normal, but to make sure, we will create visuals of every independent variable.
- linearity: the data is supposed to be linear, though we will remove some outliers if present
- Multicollinearity: as the data set is about the price of houses against the features of house, the Multicollinearity is not supposed to be present as the dependent variable follows different independent variables
- Autocorrelation: we are assuming this is not present in the data

```
[4]: data.isna().sum()
```

```
[4]: name          0
     year          0
     selling_price  0
     km_driven     0
     fuel          0
     seller_type   0
     transmission  0
     owner         0
     dtype: int64
```

```
[5]: data.corr()
```

```
[5]:          year  selling_price  km_driven
year          1.00000         0.41392    -0.41969
selling_price  0.41392         1.00000    -0.19229
km_driven     -0.41969        -0.19229     1.00000
```

```
[6]: # Making a list of all categorical variables
     cat_col = [
         "fuel",
         "seller_type",
         "transmission",
         "year",
         "owner",
```

```

]
# Printing number of count of each unique value in each column
for column in cat_col:
    print(data[column].value_counts())
    print("#" * 40)

```

```

Diesel      2153
Petrol      2123
CNG         40
LPG         23
Electric    1
Name: fuel, dtype: int64
#####
Individual  3244
Dealer      994
Trustmark Dealer  102
Name: seller_type, dtype: int64
#####
Manual      3892
Automatic   448
Name: transmission, dtype: int64
#####
2017      466
2015      421
2012      415
2013      386
2014      367
2018      366
2016      357
2011      271
2010      234
2019      195
2009      193
2008      145
2007      134
2006      110
2005       85
2020       48
2004       42
2003       23
2002       21
2001       20
1998       12
2000       12
1999       10
1997        3

```

```

1996      2
1995      1
1992      1
Name: year, dtype: int64
#####
First Owner      2832
Second Owner     1106
Third Owner       304
Fourth & Above Owner    81
Test Drive Car     17
Name: owner, dtype: int64
#####

```

2 Maximum car being sold have fuel type as Diesel.

- Mumbai has highest numbers of car available for purchase.
- Most of the cars are 5 seaters and First owned.
- Years of car ranges from 1996- 2015

```
[7]: # data processing
data.info()
```

```

<class 'pandas.core.frame.DataFrame'>
RangeIndex: 4340 entries, 0 to 4339
Data columns (total 8 columns):
#   Column          Non-Null Count  Dtype
---  -
0   name            4340 non-null   object
1   year            4340 non-null   int64
2   selling_price   4340 non-null   int64
3   km_driven       4340 non-null   int64
4   fuel            4340 non-null   object
5   seller_type     4340 non-null   object
6   transmission    4340 non-null   object
7   owner           4340 non-null   object
dtypes: int64(3), object(5)
memory usage: 271.4+ KB

```

```
[8]: data[['fuel', 'seller_type', 'transmission', 'owner']].sample(10)
```

```

[8]:      fuel seller_type transmission      owner
709   Petrol      Dealer      Manual  First Owner
2930  Petrol  Individual  Automatic  Second Owner
4316  Diesel  Individual      Manual  First Owner
1740  Petrol  Individual      Manual  Second Owner
539   Diesel  Individual  Automatic  Second Owner
3485  Petrol  Individual      Manual  First Owner

```

977	Diesel	Individual	Manual	First Owner
4296	Petrol	Dealer	Manual	First Owner
662	Petrol	Individual	Manual	First Owner
3448	Diesel	Individual	Manual	First Owner

```
[9]: data["fuel"] = data["fuel"].astype("category")
data["seller_type"] = data["seller_type"].astype("category")
data["owner"] = data["owner"].astype("category")
data["transmission"] = data["transmission"].astype("category")
```

```
[10]: data.describe().T
```

```
[10]:
```

	count	mean	std	min	25%	\
year	4340.00000	2013.09078	4.21534	1992.00000	2011.00000	
selling_price	4340.00000	504127.31175	578548.73614	20000.00000	208749.75000	
km_driven	4340.00000	66215.77742	46644.10219	1.00000	35000.00000	

	50%	75%	max
year	2014.00000	2016.00000	2020.00000
selling_price	350000.00000	600000.00000	8900000.00000
km_driven	60000.00000	90000.00000	806599.00000

```
[11]: data['Current_year']=2021
data['Ageofcar']=data['Current_year']-data['year']
data.drop('Current_year',axis=1,inplace=True)
data.head()
```

```
[11]:
```

	name	year	selling_price	km_driven	fuel	\
0	Maruti 800 AC	2007	60000	70000	Petrol	
1	Maruti Wagon R LXi Minor	2007	135000	50000	Petrol	
2	Hyundai Verna 1.6 SX	2012	600000	100000	Diesel	
3	Datsun RediGO T Option	2017	250000	46000	Petrol	
4	Honda Amaze VX i-DTEC	2014	450000	141000	Diesel	

	seller_type	transmission	owner	Ageofcar
0	Individual	Manual	First Owner	14
1	Individual	Manual	First Owner	14
2	Individual	Manual	First Owner	9
3	Individual	Manual	First Owner	4
4	Individual	Manual	Second Owner	7

```
[12]: #As mentioned in dataset car name has Brand and model so extracting it ,This
      ↪ can help to fill missing values of price column as brand
data['Brand'] = data['name'].str.split(' ').str[0] #Separating Brand name from
      ↪ the Name
data['Model'] = data['name'].str.split(' ').str[1] + data['name'].str.split('
      ↪ ').str[2]
```

```
[13]: x=data.Brand.unique()
```

```
[14]: data.Brand.nunique()
```

```
[14]: 29
```

```
[15]: y=data.groupby(data.Brand).size().sort_values(ascending =False)
```

```
[16]: data.Model.isnull().sum()
```


```
[16]: 1
```

```
[17]: data.dropna(subset=['Model'],axis=0,inplace=True)
```

```
[18]: data.Model.nunique()
```

```
[18]: 612
```

```
[19]: fig, ax = plt.subplots(figsize=(10,7))  
ax.bar(x,y, color = 'b', width = 0.8)  
plt.xticks(rotation=90)  
plt.show()
```



```
[20]: data.groupby('Model')['Model'].size().nlargest(30)
```


```
[20]: Model
      WagonR          164
      SwiftDzire      139
      Grandi10        112
      Alto800         87
      Innova2.5       83
      Verna1.6        71
      SantroXing      68
      IndicaVista     67
      AltoLXi         62
      SwiftVDI        60
      AltoK10         51
      AltoLX          47
      Creta1.6        46
      EONera          43
      BeatDiesel      40
      i10Magna        40
      800AC           37
      FigoDiesel      36
      Duster85PS      35
      EcoSport1.5     35
      i20Asta         32
      Cityi           31
      XUV500W8        31
      VernaCRDi       30
      EONMagna        29
      Spark1.0        29
      BoleroPower     28
      KWIDRXT         28
      NewSafari       28
      ZenEstilo       28
      Name: Model, dtype: int64
```

3 wagonR is the most popular model

```
[21]: plt.style.use('ggplot')
      #select all quantitative columns for checking the spread
      numeric_columns = data.select_dtypes(include=np.number).columns.tolist()
      plt.figure(figsize=(20,25))

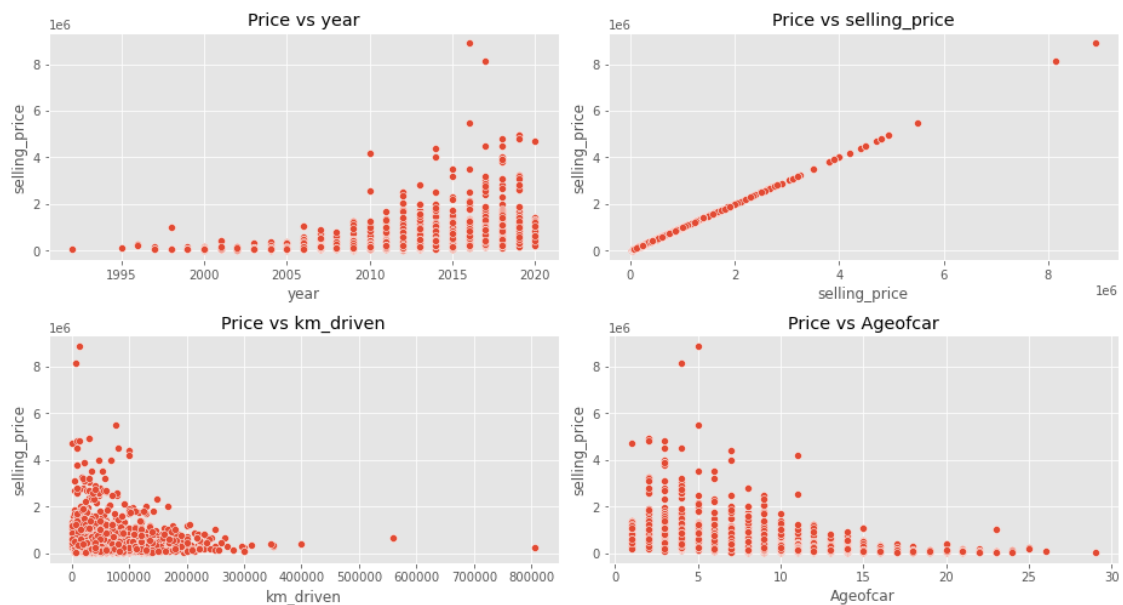
      for i, variable in enumerate(numeric_columns):
          plt.subplot(10,3,i+1)
```

```
sns.distplot(data[variable],kde=False,color='blue')
plt.tight_layout()
plt.title(variable)
```



```
[22]: numeric_columns= numeric_columns = data.select_dtypes(include=np.number).
      ↪ columns.tolist()
      plt.figure(figsize=(13,17))

      for i, variable in enumerate(numeric_columns):
          plt.subplot(5,2,i+1)
          sns.scatterplot(x=data[variable],y=data['selling_price']).
          ↪ set(title='Price vs '+ variable)
             #plt.xticks(rotation=90)
          plt.tight_layout()
```




```
[23]: # grouping the cars by high ot low profile cars
Low=['Maruti',
     'Hyundai',
     'Ambassdor',
     'Hindustan',
     'Force',
     'Chevrolet',
     'Fiat',
     'Tata',
     'Smart',
     'Renault',
     'Datsun',
     'Mahindra',
     'Skoda',
     'Ford',
     'Toyota',
     'Isuzu',
     'Mitsubishi','Honda','Land', 'Daewoo', 'MG', 'Ambassador', 'Kia',
     'OpelCorsa']
High=['Audi',
      'Mini Cooper',
      'Bentley',
      'Mercedes-Benz',
      'Lamborghini',
      'Volkswagen',
      'Porsche',
      'Land Rover',
      'Nissan',
      'Volvo',
      'Jeep',
      'Jaguar',
      'BMW']# more than 30lakh
```


```
[24]: def classrange(x):
      if x in Low:
          return "Low"
      elif x in High:
          return "High"
      else:
          return x
```

```
[25]: data['Brand_Class'] = data['Brand'].apply(lambda x: classrange(x))
```

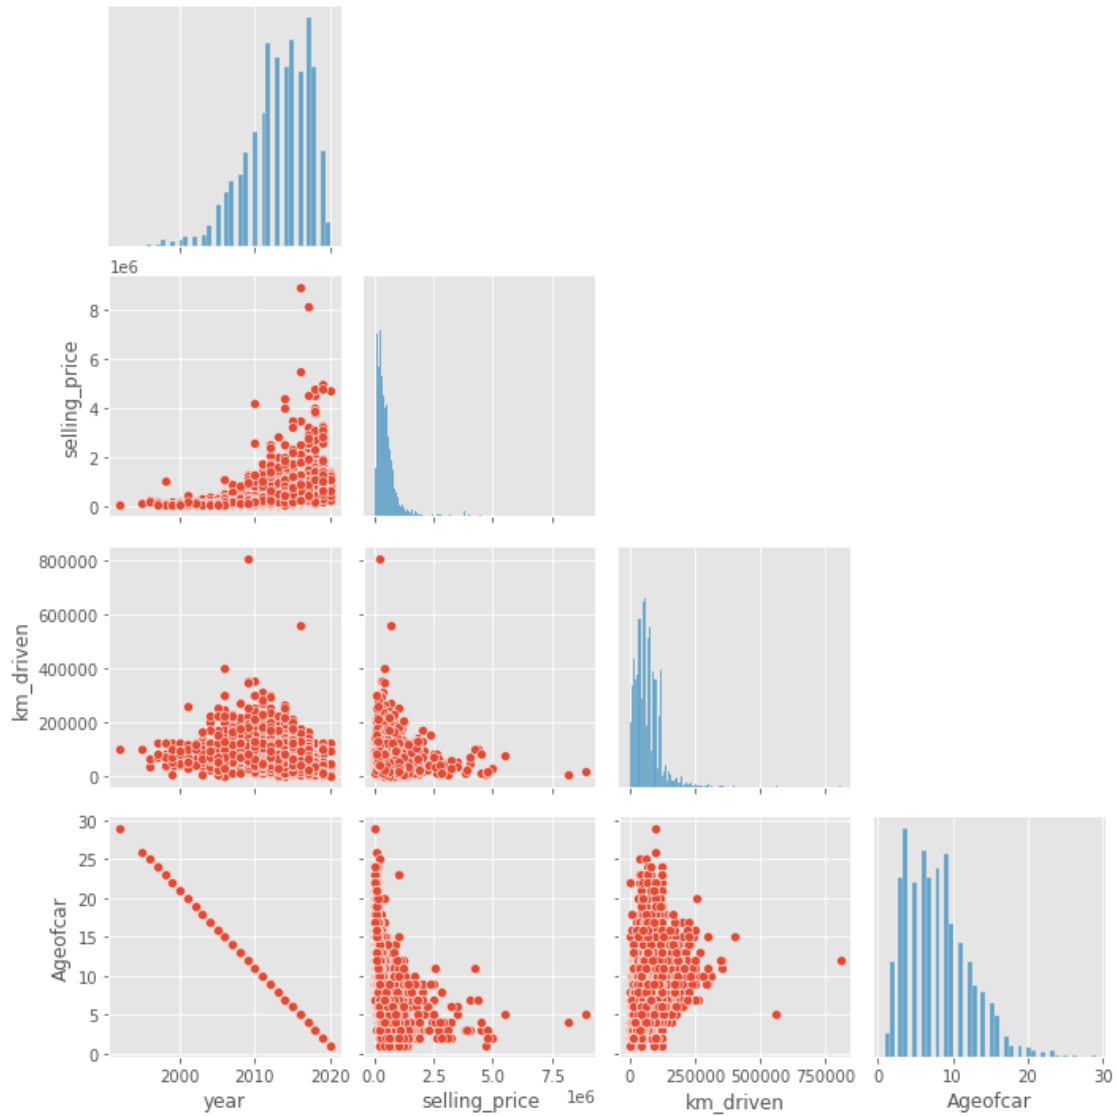
```
[26]: data['Brand_Class'].unique()
```

```
[26]: array(['Low', 'High'], dtype=object)
```

```
[27]: plt.figure(figsize=(10,8))
sns.heatmap(data.corr(),annot=True ,cmap="YlGnBu" )
plt.show()
```




```
[28]: sns.pairplot(data=data, corner=True)
plt.show()
```




```
[29]: plt.figure(figsize=(10,7))
plt.title("Price Vs Mileage")
sns.scatterplot(y='selling_price', x='km_driven', hue='fuel', data=data)
```

```
[29]: <AxesSubplot:title={'center':'Price Vs Mileage'}, xlabel='km_driven',
ylabel='selling_price'>
```

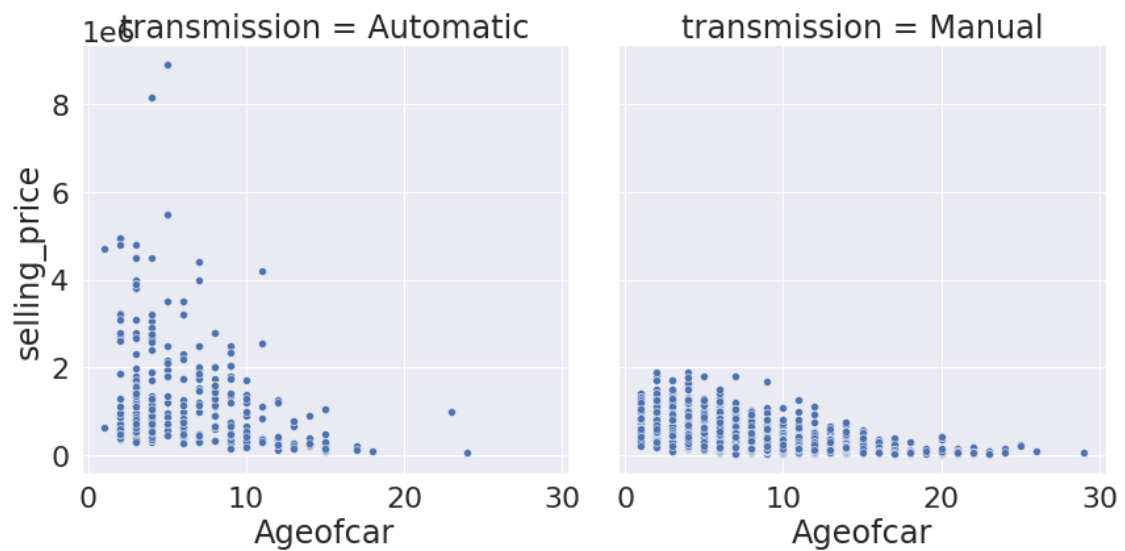


```
[30]: plt.figure(figsize=(20,15))
sns.set(font_scale=2)
sns.scatterplot(x='selling_price', y='Brand', data=data)
plt.grid()
```




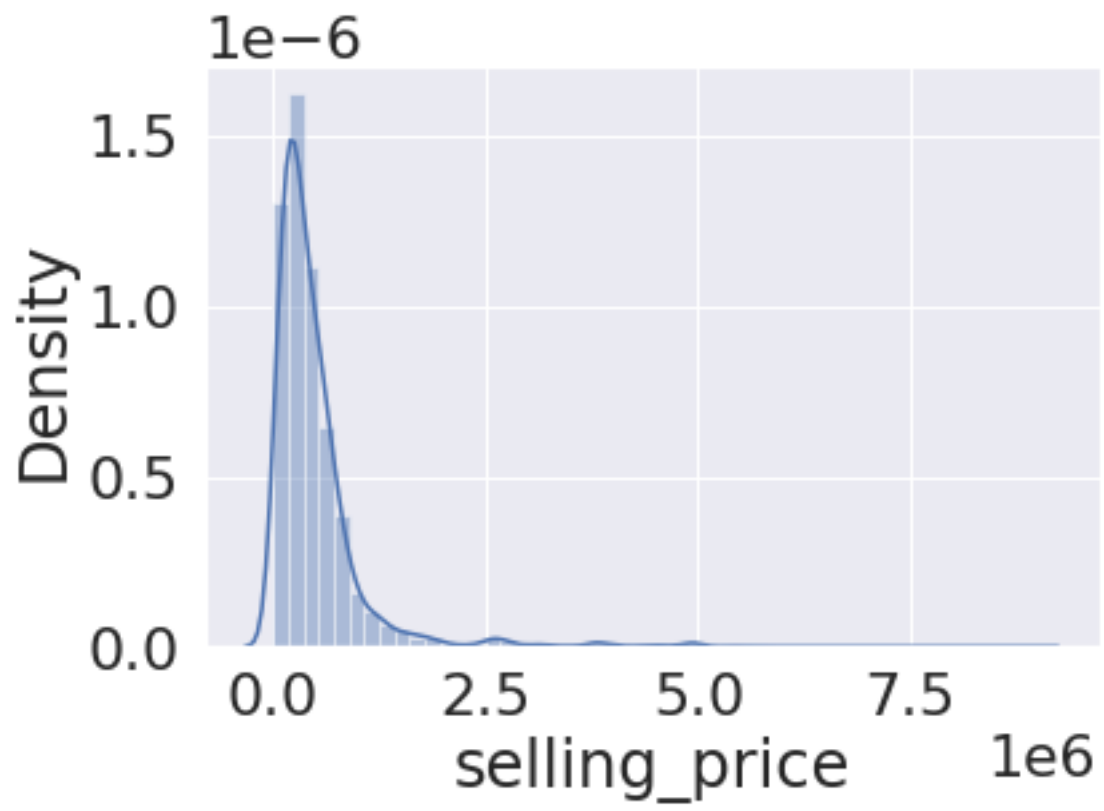
```
[31]: sns.relplot(data=data,
→y='selling_price',x='Ageofcar',col='transmission',aspect=1,height=6)
```


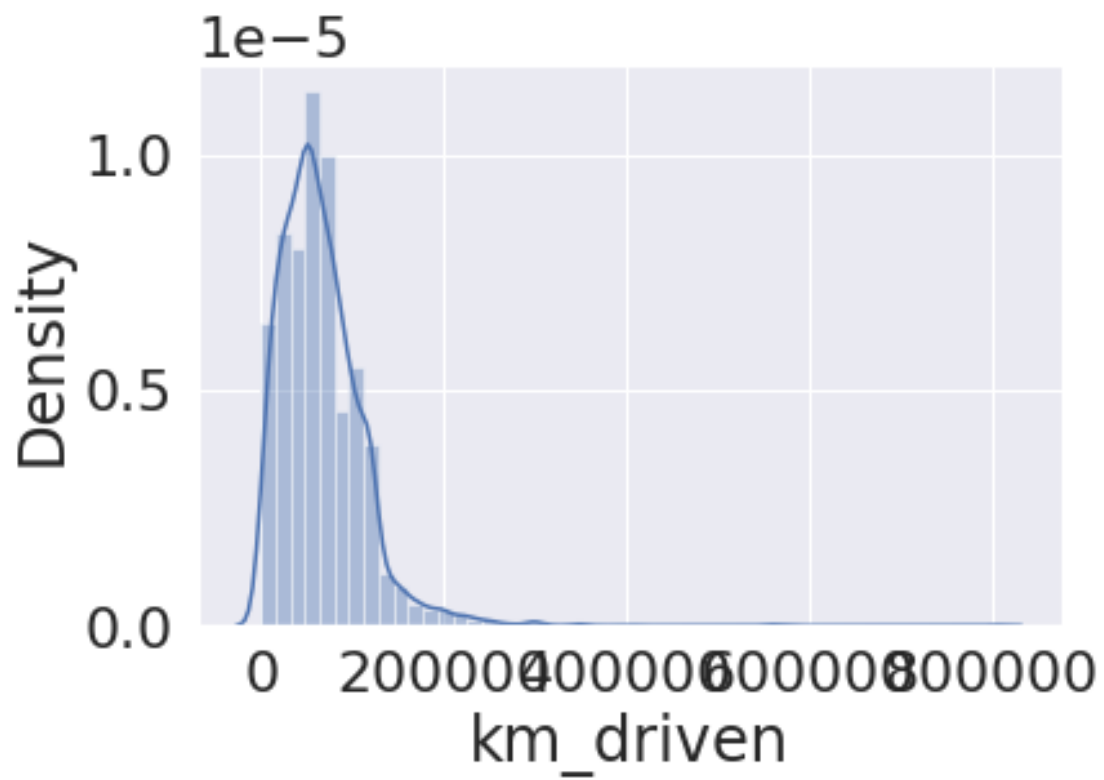
```
[31]: <seaborn.axisgrid.FacetGrid at 0x7f8d555aa6a0>
```



```
[32]: # check distrubution if skewed. If distrubution is skewed , it is advice to use ↵  
      ↪ log transform  
cols_to_log = data.select_dtypes(include=np.number).columns.tolist()  
for colname in cols_to_log:  
    sns.distplot(data[colname], kde=True)  
    plt.show()
```








```
[33]: def Perform_log_transform(df,col_log):
        """#Perform Log Transformation of dataframe , and list of columns """
        for colname in col_log:
            df[colname + '_log'] = np.log(df[colname])
        #df.drop(col_log, axis=1, inplace=True)
        df.info()
```

```
[34]: Perform_log_transform(data,['km_driven','selling_price'])
```

```
<class 'pandas.core.frame.DataFrame'>
Int64Index: 4339 entries, 0 to 4339
Data columns (total 14 columns):
#   Column                Non-Null Count  Dtype
---  -
0   name                  4339 non-null   object
1   year                  4339 non-null   int64
2   selling_price         4339 non-null   int64
3   km_driven              4339 non-null   int64
4   fuel                  4339 non-null   category
5   seller_type            4339 non-null   category
6   transmission           4339 non-null   category
7   owner                  4339 non-null   category
8   Ageofcar              4339 non-null   int64
9   Brand                  4339 non-null   object
10  Model                  4339 non-null   object
11  Brand_Class            4339 non-null   object
12  km_driven_log          4339 non-null   float64
13  selling_price_log      4339 non-null   float64
dtypes: category(4), float64(2), int64(4), object(4)
memory usage: 390.5+ KB
```

```
[35]: data.drop(['name','Model','year','Brand'],axis=1,inplace=True)
```

```
[36]: data.info()
```

```
<class 'pandas.core.frame.DataFrame'>
Int64Index: 4339 entries, 0 to 4339
Data columns (total 10 columns):
#   Column                Non-Null Count  Dtype
---  -
0   selling_price         4339 non-null   int64
1   km_driven              4339 non-null   int64
2   fuel                  4339 non-null   category
3   seller_type            4339 non-null   category
4   transmission           4339 non-null   category
5   owner                  4339 non-null   category
```

```


6   Ageofcar          4339 non-null   int64
7   Brand_Class       4339 non-null   object
8   km_driven_log     4339 non-null   float64
9   selling_price_log 4339 non-null   float64
dtypes: category(4), float64(2), int64(3), object(1)
memory usage: 254.9+ KB

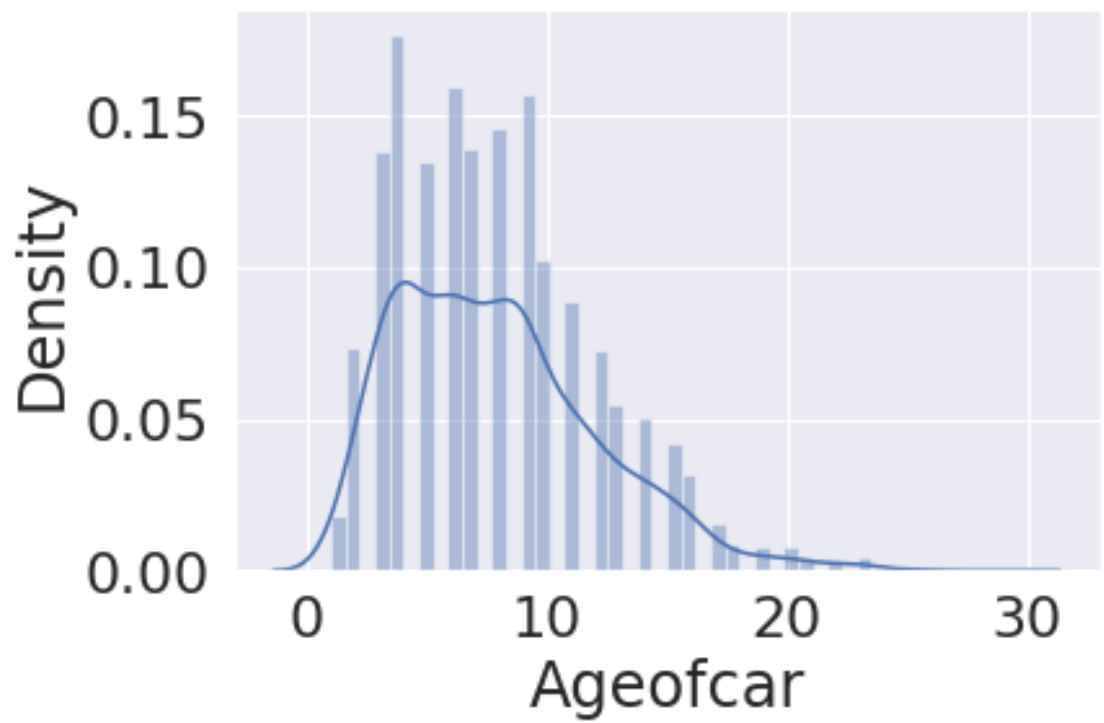

```



```

[37]: cols_to_log = data.select_dtypes(include=np.number).columns.tolist()
      for colname in cols_to_log:
          sns.distplot(data[colname], kde=True)
          plt.show()

```







4 model building

4.0.1 creating first model with only positive correlated columns

```
[38]: data.head(2)
```

```
[38]:   selling_price  km_driven   fuel seller_type transmission   owner \
0         60000      70000  Petrol  Individual         Manual  First Owner
1        135000      50000  Petrol  Individual         Manual  First Owner

   Ageofcar Brand_Class  km_driven_log  selling_price_log
0         14         Low      11.15625         11.00210
1         14         Low      10.81978         11.81303
```

```
[39]: data=data.drop(['selling_price'],axis=1)
```

```
[40]: data.head()
```

```
[40]:   km_driven   fuel seller_type transmission   owner Ageofcar \
0      70000  Petrol  Individual         Manual  First Owner      14
1      50000  Petrol  Individual         Manual  First Owner      14
2     100000  Diesel  Individual         Manual  First Owner       9
3      46000  Petrol  Individual         Manual  First Owner       4
4     141000  Diesel  Individual         Manual  Second Owner      7

   Brand_Class  km_driven_log  selling_price_log
0         Low      11.15625         11.00210
1         Low      10.81978         11.81303
2         Low      11.51293         13.30468
3         Low      10.73640         12.42922
4         Low      11.85652         13.01700
```

```
[41]: X = data.drop(["selling_price_log", 'km_driven', 'Ageofcar'], axis=1)
y = data[["selling_price_log"]]
```

```
[42]: def encode_cat_vars(x):
      x = pd.get_dummies(
          x,
          columns=x.select_dtypes(include=["object", "category"]).columns.
          ↪to_list(),
          drop_first=True,
      )
      return x
```

```
[43]: X = encode_cat_vars(X)
X.head()
```

```
[43]:   km_driven_log  fuel_Diesel  fuel_Electric  fuel_LPG  fuel_Petrol  \
0      11.15625          0          0          0          1
1      10.81978          0          0          0          1
2      11.51293          1          0          0          0
3      10.73640          0          0          0          1
4      11.85652          1          0          0          0

   seller_type_Individual  seller_type_Trustmark Dealer  transmission_Manual  \
0                      1                      0          1
1                      1                      0          1
2                      1                      0          1
3                      1                      0          1
4                      1                      0          1

   owner_Fourth & Above Owner  owner_Second Owner  owner_Test Drive Car  \
0                      0                      0          0
1                      0                      0          0
2                      0                      0          0
3                      0                      0          0
4                      0                      1          0

   owner_Third Owner  Brand_Class_Low
0          0          1
1          0          1
2          0          1
3          0          1
4          0          1
```

```
[44]: X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.3,
↳random_state=42)
X_train.reset_index()
print("X_train:",X_train.shape)
print("X_test:",X_test.shape)
print("y_train:",y_train.shape)
print("y_test:",y_test.shape)
```

```
X_train: (3037, 13)
X_test: (1302, 13)
y_train: (3037, 1)
y_test: (1302, 1)
```

```
[45]: # Statsmodel api does not add a constant by default. We need to add it
↳explicitly.
X_train = sm.add_constant(X_train)
```

```

# Add constant to test data
X_test = sm.add_constant(X_test)

def build_ols_model(train):
    # Create the model
    olsmodel = sm.OLS(y_train["selling_price_log"], train)
    return olsmodel.fit()
#fit statmodel
olsmodel1 = build_ols_model(X_train)
print(olsmodel1.summary())

```

OLS Regression Results

```

=====
Dep. Variable:      selling_price_log      R-squared:      0.513
Model:              OLS                    Adj. R-squared:  0.511
Method:             Least Squares          F-statistic:    245.4
Date:               Mon, 13 Jun 2022       Prob (F-statistic): 0.00
Time:               04:53:18               Log-Likelihood: -2693.4
No. Observations:   3037                   AIC:           5415.
Df Residuals:       3023                   BIC:           5499.
Df Model:           13
Covariance Type:    nonrobust
=====

```

	coef	std err	t	P> t
const	16.4916	0.180	91.767	0.000
km_driven_log	-0.2442	0.013	-18.258	0.000
fuel_Diesel	0.5360	0.106	5.077	0.000
fuel_Electric	-0.4174	0.600	-0.696	0.486
fuel_LPG	-0.2872	0.185	-1.557	0.120
fuel_Petrol	-0.1569	0.105	-1.487	0.137
seller_type_Individual	-0.0816	0.028	-2.959	0.003
seller_type_Trustmark Dealer	0.5276	0.074	7.140	0.000
transmission_Manual	-0.7914	0.039	-20.072	0.000

-0.869	-0.714				
owner_Fourth & Above Owner	-0.7552	0.076	-9.884	0.000	
-0.905	-0.605				
owner_Second Owner	-0.3278	0.026	-12.530	0.000	
-0.379	-0.276				
owner_Test Drive Car	0.0824	0.170	0.485	0.628	
-0.251	0.415				
owner_Third Owner	-0.5366	0.044	-12.275	0.000	
-0.622	-0.451				
Brand_Class_Low	-0.4077	0.045	-8.968	0.000	
-0.497	-0.319				

```
=====
Omnibus:                    52.817    Durbin-Watson:                1.969
Prob(Omnibus):              0.000    Jarque-Bera (JB):            56.969
Skew:                      -0.297    Prob(JB):                   4.26e-13
Kurtosis:                  3.311    Cond. No.                   620.
=====
```

Notes:

[1] Standard Errors assume that the covariance matrix of the errors is correctly specified.

5 second model

5.0.1 using all the columns

```
[46]: data.head()
```

```
[46]:   km_driven  fuel seller_type transmission      owner  Ageofcar  \
0      70000  Petrol  Individual      Manual  First Owner      14
1      50000  Petrol  Individual      Manual  First Owner      14
2     100000  Diesel  Individual      Manual  First Owner       9
3      46000  Petrol  Individual      Manual  First Owner       4
4     141000  Diesel  Individual      Manual  Second Owner       7
```

```
   Brand_Class  km_driven_log  selling_price_log
0         Low      11.15625      11.00210
1         Low      10.81978      11.81303
2         Low      11.51293      13.30468
3         Low      10.73640      12.42922
4         Low      11.85652      13.01700
```

```
[47]: X=data.drop(['selling_price_log'],axis=1)
      y=data['selling_price_log']
```

```
[48]: def encode_cat_vars(x):
      x = pd.get_dummies(
```



```

        x,
        columns=x.select_dtypes(include=["object", "category"]).columns.
        ↪tolist(),
        drop_first=True,
    )
    return x

```

```

[49]: X = encode_cat_vars(X)
      X.head()

```

```

[49]:   km_driven  Ageofcar  km_driven_log  fuel_Diesel  fuel_Electric  fuel_LPG  \
0      70000      14      11.15625          0          0          0
1      50000      14      10.81978          0          0          0
2     100000       9      11.51293          1          0          0
3      46000       4      10.73640          0          0          0
4     141000       7      11.85652          1          0          0

      fuel_Petrol  seller_type_Individual  seller_type_Trustmark Dealer  \
0              1                      1                      0
1              1                      1                      0
2              0                      1                      0
3              1                      1                      0
4              0                      1                      0

      transmission_Manual  owner_Fourth & Above Owner  owner_Second Owner  \
0                      1                      0                      0
1                      1                      0                      0
2                      1                      0                      0
3                      1                      0                      0
4                      1                      0                      1

      owner_Test Drive Car  owner_Third Owner  Brand_Class_Low
0                      0                      0              1
1                      0                      0              1
2                      0                      0              1
3                      0                      0              1
4                      0                      0              1

```

```

[50]: X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.3,
        ↪random_state=42)
      X_train.reset_index()
      print("X_train:",X_train.shape)
      print("X_test:",X_test.shape)
      print("y_train:",y_train.shape)
      print("y_test:",y_test.shape)

```

```

X_train: (3037, 15)

```

```
X_test: (1302, 15)
y_train: (3037,)
y_test: (1302,)
```

```
[51]: y_train
```

```
[51]: 929      12.34583
      2799      12.97154
      2117      12.25486
      2880      13.15192
      246       13.04979
      ...
      3445      11.00210
      466       11.79810
      3093      12.95984
      3773      12.12811
      860       12.07254
      Name: selling_price_log, Length: 3037, dtype: float64
```

```
[52]: # Statsmodel api does not add a constant by default. We need to add it
      ↪ explicitly.
      X_train = sm.add_constant(X_train)
      # Add constant to test data
      X_test = sm.add_constant(X_test)

      model2 = sm.OLS(y,X)
      model2 = model2.fit()
      model2.summary()
```

```
[52]: <class 'statsmodels.iolib.summary.Summary'>
      """
                                OLS Regression Results
      =====
      Dep. Variable:      selling_price_log    R-squared (uncentered):
      0.996
      Model:                OLS    Adj. R-squared (uncentered):
      0.996
      Method:                Least Squares    F-statistic:
      7.326e+04
      Date:                Mon, 13 Jun 2022    Prob (F-statistic):
      0.00
      Time:                04:53:19    Log-Likelihood:
      -5193.2
      No. Observations:      4339    AIC:
      1.042e+04
```

```

Df Residuals:          4324    BIC:
1.051e+04
Df Model:              15
Covariance Type:      nonrobust
=====
=====
                                coef    std err          t      P>|t|
-----
[0.025    0.975]
-----
km_driven          -1.46e-05    3.6e-07   -40.492    0.000
-1.53e-05   -1.39e-05
Ageofcar          -0.1467      0.004   -39.889    0.000
-0.154     -0.139
km_driven_log       1.1018      0.012    89.104    0.000
1.078      1.126
fuel_Diesel         3.7638      0.114    33.092    0.000
3.541      3.987
fuel_Electric       3.6455      0.812     4.489    0.000
2.053      5.238
fuel_LPG            3.0706      0.202    15.189    0.000
2.674      3.467
fuel_Petrol         3.4075      0.113    30.274    0.000
3.187      3.628
seller_type_Individual -0.1007      0.031    -3.234    0.001
-0.162     -0.040
seller_type_Trustmark Dealer 0.3937      0.084     4.687    0.000
0.229      0.558
transmission_Manual -0.5973      0.044   -13.431    0.000
-0.685     -0.510
owner_Fourth & Above Owner -0.0226      0.094    -0.240    0.811
-0.208      0.162
owner_Second Owner  -0.0682      0.031    -2.172    0.030
-0.130     -0.007
owner_Test Drive Car   3.0741      0.199    15.424    0.000
2.683      3.465
owner_Third Owner     -0.1490      0.052    -2.849    0.004
-0.252     -0.046
Brand_Class_Low       0.0142      0.051     0.279    0.780
-0.086      0.114
=====
Omnibus:           2408.720    Durbin-Watson:           1.860
Prob(Omnibus):      0.000    Jarque-Bera (JB):       37072.363
Skew:              2.314    Prob(JB):               0.00
Kurtosis:          16.551    Cond. No.               5.40e+06
=====

```

Notes:

[1] R^2 is computed without centering (uncentered) since the model does not contain a constant.

[2] Standard Errors assume that the covariance matrix of the errors is correctly specified.

[3] The condition number is large, $5.4e+06$. This might indicate that there are strong multicollinearity or other numerical problems.


"""

```
[54]: y_pred=model2.predict(X)
      y_pred
```

```
[54]: 0      11.94016
      1      11.86135
      2      12.98502
      3      13.29476
      4      12.99034
      ...
      4335    13.25623
      4336    13.25623
      4337    12.16327
      4338    13.60165
      4339    13.08165
      Length: 4339, dtype: float64
```

```
[55]: sns.regplot(x=y, y=y_pred, ci=None, color="b")
```

```
[55]: <AxesSubplot:xlabel='selling_price_log'>
```



6 Recommendations and insights

- need to acquire more automatic cars
- should focus on new cars to gain more profit
- diesel cars are more popular
- first degree of owner has more value as price
- budget cars are popular
- high profile cars have low significant values

[]: